

HW1 Associative Model Forecast

Group 2

2025-09-13

Table of contents

| | | |
|---|--------------|----|
| 1 | Explain Data | 1 |
| 2 | Test | 3 |
| 3 | ARIMAX model | 4 |
| 4 | Forecast | 5 |
| 5 | Forecast 2 | 8 |
| 6 | Plot | 12 |
| 7 | Conclusion | 15 |

中文

1 Explain Data

已對資料做整理，取每個月的最後一筆收盤價當成月的基礎。

再利用轉對數的形式讓小樣本的序列能更平穩

報酬率公式:

$$r_t = \ln \left(\frac{P_t}{P_{t-1}} \right)$$

利用報酬率來預測輝達收盤價跟SOXX(美國半導體產業股票指數)的收盤價的關聯

本次使用ARIMAX 模型來探討，但因資料筆數為11筆，模型只能作為參考。

```
library(readr)
library(forecast)
```

```
Registered S3 method overwritten by 'quantmod':
  method          from
as.zoo.data.frame zoo
```

```
library(skimr)
library(knitr)
library(kableExtra)
data <- read_csv("NVDA_SOXX_monthly_prices_and_logreturns.csv")
```

Rows: 11 Columns: 5

```
-- Column specification -----
Delimiter: ","
chr (1): MonthEnd
dbl (4): NVDA, SOXX, NVDA_ret, SOXX_ret
```

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
skim(data)
```





Table 1: Data summary

| | |
|------------------------|------|
| Name | data |
| Number of rows | 11 |
| Number of columns | 5 |
| Column type frequency: | |
| character | 1 |
| numeric | 4 |
| Group variables | None |

Variable type: character

| skim_variable | n_missing | complete_rate | min | max | empty | n_unique | whitespace |
|---------------|-----------|---------------|-----|-----|-------|----------|------------|
| MonthEnd | 0 | 1 | 9 | 10 | 0 | 11 | 0 |

Variable type: numeric

| skim_variable | n_missing | complete_rate | mean | sd | p0 | p25 | p50 | p75 | p100 | hist |
|---------------|-----------|---------------|--------|-------|--------|--------|--------|--------|--------|---|
| NVDA | 0 | 1 | 137.52 | 23.58 | 108.38 | 122.50 | 134.29 | 148.12 | 177.87 |  |
| SOXX | 0 | 1 | 216.07 | 19.87 | 183.84 | 206.73 | 215.49 | 228.48 | 245.32 |  |
| NVDA_ret | 0 | 1 | 0.03 | 0.11 | - | - | 0.04 | 0.10 | 0.22 |  |
| | | | | | 0.14 | 0.03 | | | | |
| SOXX_ret | 0 | 1 | 0.01 | 0.07 | - | - | 0.00 | 0.02 | 0.15 |  |
| | | | | | 0.10 | 0.03 | | | | |

2 Test

下面使用adf test 來檢驗，因資料筆數的不夠，
導致p-value值過高或出現NA

```
library(tseries)
library(knitr)
library(broom)
nvda_adf <- adf.test(data$NVDA_ret)
soxx_adf <- adf.test(data$SOXX_ret)
kable(tidy(nvda_adf), caption = "ADF test on NVDA returns", digits = 4)
```

Table 4: ADF test on NVDA returns

| statistic | p.value | parameter | method | alternative |
|-----------|---------|-----------|------------------------------|-------------|
| -1.8792 | 0.6184 | 2 | Augmented Dickey-Fuller Test | stationary |

```
kable(tidy(soxx_adf), caption = "ADF test on SOXX returns", digits = 4)
```

Table 5: ADF test on SOXX returns

| statistic | p.value | parameter | method | alternative |
|-----------|---------|-----------|------------------------------|-------------|
| -0.8332 | 0.9451 | 2 | Augmented Dickey-Fuller Test | stationary |

3 ARIMAX model

下面讓我們的輝達股價為 Y ，讓SOXX的股價為 X ，

由於ARIMA(0,0,0)，所以在ARIMAX裡只做簡單線性回歸，因資料筆數太少

探討是否因為SOXX的股票升漲，進而影響輝達的股價，

可以看出在 ARIMAX 模型中，SOXX 的迴歸係數估計為 1.0253 (標準誤 0.3631，t 值 \square 2.82，顯著)，

這表示在控制自相關結構後，SOXX 報酬與 NVDA 報酬之間存在顯著的正向關聯：

SOXX 每上漲 1%，NVDA 約上漲 1.025%。

這符合金融直覺，因為 SOXX 代表整體半導體產業，而 NVDA 作為產業龍頭

，其股價自然與產業走勢緊密連動。

誤差項可以看出ACF1：殘差一階自相關 \square -0.19，代表殘差大致沒有強烈自相關。

百分比誤差超過150%應該是因為 y_{ts} 值太接近零 (導致分母小)。

```
y <- data$NVDA_ret
X <- stats::lag(data$SOXX_ret, -1) # 1 t NVDA t-1 SOXX
df <- data.frame(y=y, X=X)
df <- na.omit(df)
```

```
y_ts <- ts(df$y, frequency = 12, start = c(2024,8)) # 2024 8
X_mat <- as.matrix(df$X)
# ---- ARIMAX ----
fit <- auto.arima(y_ts, xreg = X_mat, seasonal = FALSE,
                  stepwise = FALSE, approximation = FALSE)

library(broom)
library(knitr)
library(kableExtra)
```

```
#
coef_tab <- tidy(fit)
kable(coef_tab, format="latex", booktabs=TRUE,
      caption="ARIMAX Model Coefficients") %>%
  kable_styling(latex_options="scale_down", position="center")
```

```
# (AIC, BIC, logLik)
model_info <- glance(fit)
kable(model_info, format="latex", booktabs=TRUE,
      caption="ARIMAX Model Fit Summary") %>%
  kable_styling(latex_options="scale_down", position="center")
```

Table 6: ARIMAX Model Coefficients

| term | estimate | std.error |
|------|----------|-----------|
| xreg | 1.025308 | 0.3630819 |

Table 7: ARIMAX Model Fit Summary

| sigma | logLik | AIC | BIC | nobs |
|-----------|---------|-----------|-----------|------|
| 0.0868561 | 11.7944 | -19.58881 | -18.79302 | 11 |

```
bp <- tidy(Box.test(fit$residuals, type="Box-Pierce", lag=1))
kable(bp, format="latex", booktabs=TRUE,
      caption="Box-Pierce Test on Residuals") %>%
  kable_styling(latex_options="scale_down", position="center")
```

4 Forecast

下面用報酬預測9~12月的預測再轉為收盤價得出結果，使用了天真法預測。

```
future_X <- matrix(rep(tail(df$X,1), 4), ncol = 1)
colnames(future_X) <- "X"
fc <- forecast(fit, xreg = future_X, h=4)
```

Warning in forecast.forecast_ARIMA(fit, xreg = future_X, h = 4): xreg contains different column names from the xreg used in training. Please check that the regressors are in the same order.

```
library(knitr)
library(kableExtra)

fc_table <- data.frame(
  Step = 1:length(fc$mean),
```

Table 8: Box-Pierce Test on Residuals

| statistic | p.value | parameter | method |
|-----------|-----------|-----------|-----------------|
| 0.3991365 | 0.5275356 | 1 | Box-Pierce test |

Table 9: NVDA Forecast (h = 4)

| Step | Mean | Lo80 | Hi80 | Lo95 | Hi95 |
|------|----------|-----------|----------|-----------|----------|
| 1 | 0.022351 | -0.088959 | 0.133662 | -0.147884 | 0.192586 |
| 2 | 0.022351 | -0.088959 | 0.133662 | -0.147884 | 0.192586 |
| 3 | 0.022351 | -0.088959 | 0.133662 | -0.147884 | 0.192586 |
| 4 | 0.022351 | -0.088959 | 0.133662 | -0.147884 | 0.192586 |

```

Mean = round(as.numeric(fc$mean), 6),
Lo80 = round(as.numeric(fc$lower[, "80%"]), 6),
Hi80 = round(as.numeric(fc$upper[, "80%"]), 6),
Lo95 = round(as.numeric(fc$lower[, "95%"]), 6),
Hi95 = round(as.numeric(fc$upper[, "95%"]), 6)
)

kable(
  fc_table,
  format = "latex", booktabs = TRUE,
  caption = "NVDA Forecast (h = 4)",
  align = c("c", rep("r", 5))
) %>%
  kable_styling(latex_options = "scale_down", position = "center")

```

```

h <- 4
ret4 <- as.numeric(tail(fc$mean, h))

# →
last_price <- as.numeric(tail(data$NVDA, 1))
price4 <- last_price * cumprod(exp(ret4))

# =4
pred_table <- data.frame(
  Month = c("2025-09", "2025-10", "2025-11", "2025-12"),
  NVDA_ret_forecast = ret4,
  NVDA_price_forecast = price4,
  stringsAsFactors = FALSE
)

library(knitr)
library(kableExtra)

```

| Month | NVDA_ret_forecast | NVDA_price_forecast |
|---------|-------------------|---------------------|
| 2025-09 | 0.0223513 | 178.1170 |
| 2025-10 | 0.0223513 | 182.1429 |
| 2025-11 | 0.0223513 | 186.2599 |
| 2025-12 | 0.0223513 | 190.4699 |

```
library(dplyr)
```

Attaching package: 'dplyr'

The following object is masked from 'package:kableExtra':

group_rows

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
kable(
  pred_table %>%
    mutate(
      NVDA_ret_forecast = round(NVDA_ret_forecast, 8),
      NVDA_price_forecast = round(NVDA_price_forecast, 4)
    ),
  format = "latex", booktabs = TRUE,
  col.names = c("Month", "NVDA_ret_forecast", "NVDA_price_forecast"),
  align = c("c", "r", "r")
) %>%
  kable_styling(latex_options = "scale_down", position = "center")
```

5 Forecast 2

這段使用了蒙地卡羅的方法預測，展開了不確定性，模擬多條SOXX的可能路徑，處理了未來外生變數在傳遞到NVDA的ARIMAX。

我們先對SOXX進行ARIMA預估，再丟入NVDA的ARIMAX進行股價的預測。

並產生三個悲觀和正常和樂觀的三個部分

```
set.seed(42)      #
h <- 4            # NVDA 4
# NVDA_t SOXX_{t-1} 3 10~12

y <- data$NVDA_ret
X <- stats::lag(data$SOXX_ret, -1)      # t NVDA t-1 SOXX
df <- na.omit(data.frame(y = y, X = X))

y_ts <- ts(df$y, frequency = 12, start = c(2024, 8))
X_mat <- as.matrix(df$X)

fit_nvda <- auto.arima(
  y_ts, xreg = X_mat, seasonal = FALSE,
  stepwise = FALSE, approximation = FALSE
)

# (1) SOXX_ret ARIMA SOXX

soxx_ts <- ts(na.omit(data$SOXX_ret), frequency = 12, start = c(2024, 8))
fit_soxx <- auto.arima(
  soxx_ts, seasonal = FALSE,
  stepwise = FALSE, approximation = FALSE
)

# (2) SOXX_ret 3
# simulate 2000 times

n_sim <- 2000
lastX <- as.numeric(tail(soxx_ts, 1))

# 3 SOXX ARIMA 3
# SOXX 3
```



```

sim_one_soxx <- function() {
  as.numeric(simulate(fit_soxx, nsim = 3, future = TRUE))
}

soxx_future_mat <- t(replicate(n_sim, sim_one_soxx())) # n_sim x 3

# (3)      NVDA  ARIMAX
#
#      N()
#
## 95% = mean ± z * sd → sd (upper-mean)/z
draw_one_from_fc <- function(fc_obj) {
  m <- as.numeric(fc_obj$mean)
  lo <- as.numeric(fc_obj$lower[, "95%"])
  hi <- as.numeric(fc_obj$upper[, "95%"])
  z <- qnorm(0.95)
  sd <- (hi - m) / z
  rnorm(length(m), mean = m, sd = sd)
}

nvda_ret_sims <- matrix(NA_real_, nrow = n_sim, ncol = h)

for (i in 1:n_sim) {
  # lag 1 NVDA lastX 3
  future_X <- c(lastX, soxx_future_mat[i, ])
  fc_i <- forecast(fit_nvda, xreg = matrix(future_X, ncol = 1), h = h)
  #      ARIMAX
  nvda_ret_sims[i, ] <- draw_one_from_fc(fc_i)
}

# (4) →

last_price <- as.numeric(tail(data$NVDA, 1)) #
#
nvda_price_sims <- t(apply(nvda_ret_sims, 1, function(r) last_price * exp(cumsum(r))))

# (5)      P10 P90

quant <- function(x, p) as.numeric(quantile(x, probs = p, na.rm = TRUE))
months <- c("2025-09", "2025-10", "2025-11", "2025-12")
library(knitr)

```

Table 10: NVDA Forecast (Monte Carlo Simulation, 2000 runs)

| Month | Ret_Mean | Ret_Median | Ret_P10 | Ret_P90 | Price_Mean | Price_Median | Price_P10 | Price_P90 |
|---------|-----------|------------|-----------|----------|------------|--------------|-----------|-----------|
| 2025-09 | 0.022881 | 0.025947 | -0.113550 | 0.157585 | 179.20 | 178.76 | 155.48 | 203.91 |
| 2025-10 | 0.048667 | 0.051294 | -0.106559 | 0.195001 | 189.42 | 186.88 | 153.20 | 228.09 |
| 2025-11 | -0.008766 | -0.006962 | -0.171210 | 0.149056 | 189.69 | 184.97 | 142.65 | 242.42 |
| 2025-12 | -0.006890 | -0.007120 | -0.171367 | 0.164317 | 190.42 | 184.23 | 134.30 | 257.02 |

```
library(kableExtra)

summary_table <- data.frame(
  Month = months,
  Ret_Mean = apply(nvda_ret_sims, 2, mean),
  Ret_Median = apply(nvda_ret_sims, 2, median),
  Ret_P10 = apply(nvda_ret_sims, 2, quant, p = 0.10),
  Ret_P90 = apply(nvda_ret_sims, 2, quant, p = 0.90),
  Price_Mean = apply(nvda_price_sims, 2, mean),
  Price_Median = apply(nvda_price_sims, 2, median),
  Price_P10 = apply(nvda_price_sims, 2, quant, p = 0.10),
  Price_P90 = apply(nvda_price_sims, 2, quant, p = 0.90)
)

kable(
  summary_table %>%
    mutate(
      Ret_Mean = round(Ret_Mean, 6),
      Ret_Median = round(Ret_Median, 6),
      Ret_P10 = round(Ret_P10, 6),
      Ret_P90 = round(Ret_P90, 6),
      Price_Mean = round(Price_Mean, 2),
      Price_Median = round(Price_Median, 2),
      Price_P10 = round(Price_P10, 2),
      Price_P90 = round(Price_P90, 2)
    ),
  format = "latex", booktabs = TRUE,
  caption = "NVDA Forecast (Monte Carlo Simulation, 2000 runs)",
  align = c("c", rep("r", 8))
) %>%
  kable_styling(latex_options = "scale_down", position = "center")
```

```
# -----
# 6)      bear/base/bull
#      SOXX_ret  ARIMA      ±
# -----
```

```

# SOXX 3 95% → sd
fc_soxx <- forecast(fit_soxx, h = 3)
m_sx <- as.numeric(fc_soxx$mean)
sd_sx <- (as.numeric(fc_soxx$upper[, "95%"]) - m_sx) / qnorm(0.95)

scenarios <- list(
  bear = m_sx - sd_sx,
  base = m_sx,
  bull = m_sx + sd_sx
)

run_scn <- function(path3) {
  future_X <- c(lastX, path3)
  fc <- forecast(fit_nvda, xreg = matrix(future_X, ncol = 1), h = h)
  #
  r_mean <- as.numeric(fc$mean)
  price_path <- last_price * exp(cumsum(r_mean))
  data.frame(
    Month = months,
    NVDA_ret = r_mean,
    NVDA_price = price_path
  )
}

scn_bear <- run_scn(scenarios$bear)
scn_base <- run_scn(scenarios$base)
scn_bull <- run_scn(scenarios$bull)

print_scn <- function(df, title){
  kable(
    df %>%
      mutate(
        NVDA_ret = round(NVDA_ret, 6),
        NVDA_price = round(NVDA_price, 2)
      ),
    format = "latex", booktabs = TRUE,
    caption = paste("Scenario:", title),
    col.names = c("Month", "NVDA_ret_forecast", "NVDA_price_forecast"),
    align = c("c", "r", "r")
  ) %>%
  kable_styling(latex_options = "scale_down", position = "center")
}

print_scn(scn_bear, "Bear")

```

Table 11: Scenario: Bear

| Month | NVDA_ret_forecast | NVDA_price_forecast |
|---------|-------------------|---------------------|
| 2025-09 | 0.022351 | 178.12 |
| 2025-10 | -0.016030 | 175.28 |
| 2025-11 | -0.087983 | 160.52 |
| 2025-12 | -0.087983 | 147.00 |

Table 12: Scenario: Base

| Month | NVDA_ret_forecast | NVDA_price_forecast |
|---------|-------------------|---------------------|
| 2025-09 | 0.022351 | 178.12 |
| 2025-10 | 0.053336 | 187.88 |
| 2025-11 | 0.000000 | 187.88 |
| 2025-12 | 0.000000 | 187.88 |

```
print_scn(scn_base, "Base")
```

```
print_scn(scn_bull, "Bull")
```

6 Plot

```
library(dplyr)
library(ggplot2)
library(scales)
```

Attaching package: 'scales'

Table 13: Scenario: Bull

| Month | NVDA_ret_forecast | NVDA_price_forecast |
|---------|-------------------|---------------------|
| 2025-09 | 0.022351 | 178.12 |
| 2025-10 | 0.122702 | 201.37 |
| 2025-11 | 0.087983 | 219.89 |
| 2025-12 | 0.087983 | 240.11 |

The following object is masked from 'package:readr':

col_factor

```
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

date, intersect, setdiff, union

```
hist_df <- data %>%
  mutate(MonthEnd = as.Date(MonthEnd)) %>%
  filter(MonthEnd >= as.Date("2024-08-01"), MonthEnd <= as.Date("2025-08-31")) %>%
  select(MonthEnd, NVDA)

ret_mean <- as.numeric(fc$mean)
ret_lo80 <- as.numeric(fc$lower[, "80%"])
ret_hi80 <- as.numeric(fc$upper[, "80%"])
ret_lo95 <- as.numeric(fc$lower[, "95%"])
ret_hi95 <- as.numeric(fc$upper[, "95%"])

#
last_price <- as.numeric(tail(data$NVDA, 1))

px_mean <- last_price * cumprod(exp(ret_mean))
px_lo80 <- last_price * cumprod(exp(ret_lo80))
px_hi80 <- last_price * cumprod(exp(ret_hi80))
px_lo95 <- last_price * cumprod(exp(ret_lo95))
px_hi95 <- last_price * cumprod(exp(ret_hi95))

#          2025/08 →
start_month <- floor_date(max(as.Date(data$MonthEnd)), "month") + months(1)
future_months <- seq.Date(from = start_month, by = "month", length.out = length(ret_mean))

fc_df <- data.frame(
  MonthEnd = future_months,
  px_mean  = px_mean,
  px_lo80  = px_lo80,
  px_hi80  = px_hi80,
  px_lo95  = px_lo95,
```


```

    px_hi95 = px_hi95
  )

# ==== 2)      +      +      ====
g <- ggplot() +
  #
  geom_line(data = hist_df, aes(MonthEnd, NVDA),
            linewidth = 1, color = "#1f78b4") +
  # 95%
  geom_ribbon(
    data = fc_df,
    aes(x = MonthEnd, ymin = px_lo95, ymax = px_hi95),
    fill = "#fb9a99", alpha = 0.25
  ) +
  # 80%
  geom_ribbon(
    data = fc_df,
    aes(x = MonthEnd, ymin = px_lo80, ymax = px_hi80),
    fill = "#fdbf6f", alpha = 0.35
  ) +
  #
  geom_line(
    data = fc_df,
    aes(MonthEnd, px_mean),
    linetype = "dashed", linewidth = 1, color = "#e31a1c"
  ) +
  scale_x_date(
    date_labels = "%Y-%m",
    breaks = "1 month",
    limits = c(as.Date("2024-08-01"), as.Date("2025-12-31")),
    expand = expansion(mult = c(0.01, 0.02))
  ) +
  scale_y_continuous(labels = label_number(big.mark = ",")) +
  labs(
    title = "NVDA      ",
    x = "Month", y = "Price",
    caption = "          / 80%/95%  "
  ) +
  theme_minimal(base_size = 13) +
  theme(
    plot.title = element_text(face = "bold", color = "#333333"),
    axis.text.x = element_text(angle = 45, hjust = 1)
  )

```

```
print(g)
```



7 Conclusion

由於資料只有近一年的，導致ARIMAX的參數估計不是很穩定，因此分析結果只能當參考之用，上面提到的ARIMA(0,0,0)為簡單回歸，預測輝達股價會因為SOXX的股價影響， β 為 1.025308 (正相關) 可以得知當SOXX股票漲，輝達就會漲。